Numerical Solution of a Variable-Diffusion Elliptic BVP

Project 2 (PDE BVP) Report — Team 15

Chuyang Wang, 2363591 (solo)

Abstract

Steady heat conduction in a heterogeneous medium is modelled by the divergence-form elliptic problem $-\nabla \cdot (a\nabla u) = f$ on the unit square (Problem Pack Direction 3). I discretise with a conservative face-averaged five-point finite-difference scheme, assemble in COO, solve in CSR, and move Dirichlet data into the right-hand side so that boundary nodes are never unknowns. On the smooth manufactured benchmark $a = 1 + x$, $u^* = \sin(\pi x)\sin(\pi y)$, discrete L^2 and L^∞ errors both fit order 2.00 over $n = 15, 31, 63, 127$, with $\|A - A^\top\|_{\max}/\|A\|_{\max} = 0$. On the layered jump $a = 1\,|\,10$ at $x = 1/2$, a harmonic face average is nodally exact ($L^\infty \sim 10^{-14}$) and recovers the interface flux $20/11$; the arithmetic average of the same nodal values yields first-order L^∞ error and an $\mathcal{O}(1)$ flux jump under refinement. Unpreconditioned CG matches the direct discrete L^2 error once the relative residual reaches 10^{-4}; at $n = 31$ there is no runtime crossover with `spsolve`. The scheme is second-order for smooth a; across an unresolved jump an arithmetic face coefficient loses both order and flux continuity, and CG iteration counts grow with contrast a_R/a_L.

Keywords: elliptic BVP, variable diffusion, five-point stencil, harmonic average, method of manufactured solutions, sparse direct / CG

Contents

1 Introduction and Motivation **3**
- 1.1 Problem Background . . . 3
- 1.2 Mathematical Model . . . 3
- 1.3 Overview of the Report . . . 3

2 Discretisation and Analysis **4**
- 2.1 Discrete Scheme . . . 4
- 2.2 Consistency and Truncation Error . . . 4
- 2.3 Boundary Treatment . . . 5
- 2.4 Discrete Structure, Stability, and Convergence . . . 5

3 Implementation **6**
- 3.1 Algorithm . . . 6
- 3.2 Code Structure . . . 6
- 3.3 Key Design Decisions . . . 6

4 Verification and Numerical Experiments **7**
- 4.1 Verification: Method of Manufactured Solutions . . . 7
- 4.2 Separating Algebraic and Discretisation Error . . . 9
- 4.3 Depth Comparison: Harmonic versus Arithmetic Face Coefficients . . . 10
- 4.4 Main Results: Contrast, Conditioning, Cost . . . 11
- 4.5 Limitations and Robustness . . . 12

5 Conclusions **12**

A AI Transparency Log **13**

B Individual Contribution Statement **13**

C Supplementary data **14**

1 Introduction and Motivation

1.1 Problem Background

Steady heat conduction in a layered composite, soil profile, or tissue boundary is the model problem in which the conductivity $a(x, y)$ varies in space and may jump [8, 1]. The physically correct operator is the divergence form $-\nabla \cdot (a\nabla u)$: the normal flux $a\partial_n u$ is continuous across an interface. A pointwise expansion $a\Delta u + \nabla a \cdot \nabla u$ is the wrong starting point once a is discontinuous—it loses conservation, symmetry, and the M-matrix sign pattern.

I chose Direction 3 because the unit square keeps the five-point machinery of Tutorial 5 intact, while the jump forces a modelling choice (arithmetic versus harmonic face coefficients) that a smooth manufactured Poisson problem never sees.

1.2 Mathematical Model

On $\Omega = (0, 1)^2$ with Dirichlet data on $\partial\Omega$,

$$-\nabla \cdot \big(a(x, y)\nabla u\big) = f \quad \text{in } \Omega, \qquad u = g \quad \text{on } \partial\Omega, \tag{1}$$

with $0 < a_{\min} \le a(x, y) \le a_{\max} < \infty$. The sign convention makes $-\nabla \cdot (a\nabla u)$ a positive operator, so the discrete matrix has a positive diagonal.

Smooth manufactured benchmark (course bank problem P4).

$$a(x, y) = 1 + x, \qquad u^*(x, y) = \sin(\pi x)\sin(\pi y), \qquad g = 0,$$

and

$$f = 2\pi^2(1 + x)\sin(\pi x)\sin(\pi y) - \pi\cos(\pi x)\sin(\pi y). \tag{2}$$

The extra $-\pi\cos(\pi x)\sin(\pi y)$ term is $\nabla a \cdot \nabla u^*$; a constant-coefficient five-point stencil cannot capture it.

Layered-interface benchmark.

$$a(x) = \begin{cases} 1, & x < 1/2, \\ 10, & x > 1/2, \end{cases} \qquad f = 0,$$

and $g = u^*$ on $\partial\Omega$, where the exact one-dimensional profile

$$u^*(x) = \begin{cases} \dfrac{20}{11}\,x, & 0 \le x \le 1/2, \\ \dfrac{2}{11}\,x + \dfrac{9}{11}, & 1/2 \le x \le 1, \end{cases} \tag{3}$$

is independent of y. It is continuous, $u^*(0) = 0$, $u^*(1) = 1$, the derivative jumps, and the flux $au^*_x = 20/11$ is continuous.

Expected signs: on the layered problem, $u^* \in [0, 1]$ with extrema on the boundary (maximum principle). On P4, $u^* > 0$ in Ω and $u^* = 0$ on $\partial\Omega$.

1.3 Overview of the Report

Section 2 derives the face-averaged stencil and Dirichlet shift, states truncation assumptions, and links energy stability to convergence—including the series-resistance derivation of the harmonic face mean. Section 3 records the implementation. Section 4 verifies P4 by MMS, separates algebraic from

discretisation error, falsifies "symmetry ⇒ correct interface," and reports cost at a matched accuracy target. Section 5 summarises the findings.

2 Discretisation and Analysis

2.1 Discrete Scheme

I use a uniform nodal grid with n interior nodes per direction,

$$h = \frac{1}{n+1}, \qquad x_i = (i+1)h, \qquad y_j = (j+1)h, \qquad i, j = 0, \ldots, n-1. \tag{4}$$

The unknown vector is indexed by $k = in + j$ (`meshgrid` with `indexing='ij'`; axis 0 is x). Boundary nodes are not unknowns.

A conservative finite-volume reading of $-\nabla \cdot (a\nabla u)$ on the dual cell of width h gives the five-point divergence-form stencil [2, 8]

$$(A_h U)_{ij} = \frac{1}{h^2}\Big[a_E(U_{ij} - U_{i+1,j}) + a_W(U_{ij} - U_{i-1,j}) + a_N(U_{ij} - U_{i,j+1}) + a_S(U_{ij} - U_{i,j-1})\Big]. \tag{5}$$

When $a \equiv 1$ this recovers the classical five-point Laplacian

$$(-\Delta_h u)_{ij} = \frac{4u_{ij} - u_W - u_E - u_S - u_N}{h^2}.$$

The four face coefficients are shared: the east coefficient of node (i, j) is the west coefficient of node $(i+1, j)$. That sharing makes A_h symmetric after Dirichlet elimination.

For smooth a I evaluate the coefficient at the half-point, $a_E = a(x_i + h/2, y_j)$. For $a = 1 + x$ this coincides with the arithmetic nodal midpoint and is second-order. At a material interface I use the resistance (harmonic) average of the two nodal values,

$$a_E = \frac{2a_L a_R}{a_L + a_R}, \tag{6}$$

and, as the depth comparison, the arithmetic mean $(a_L + a_R)/2$. When the interface bisects a grid edge, (6) is exact for piecewise-constant a. I therefore take odd interval counts $N = n + 1 \in \{31, 63, 127, 255\}$ on the layered test, so $x = 1/2$ lies midway between two interior columns [8].

2.2 Consistency and Truncation Error

For a single east face with smooth a and u, the centred flux

$$a_E \frac{U_{i+1,j} - U_{ij}}{h} = a\big(x_i + \tfrac{h}{2}, y_j\big)\; u_x\big(x_i + \tfrac{h}{2}, y_j\big) + \mathcal{O}(h^2)$$

is second-order; summing four such faces and dividing by h recovers $-\nabla \cdot (a\nabla u) + \mathcal{O}(h^2)$. The same u_{xxxx} and u_{yyyy} pieces appear as in the five-point Laplacian, plus smooth lower-order terms from ∇a. Assumptions: $u \in C^4(\overline{\Omega})$ and $a \in C^2(\overline{\Omega})$. Across the layered jump, u^* is only Lipschitz and $a \notin C^2$, so this expansion fails at the interface face; second-order truncation is not claimed for an arithmetic coefficient there.

I measure the truncation residual $\tau_h = b - A_h u^*$ in the discrete L^2 norm $h\|\tau_h\|_2$. On P4 this quantity drops by four under each mesh doubling (Table 2), consistent with $\mathcal{O}(h^2)$ pointwise truncation. On the four meshes of Table 1 there is no finest-mesh floor: both named errors continue to quarter, so the fitted order 2.00 is not contaminated by round-off.

2.3 Boundary Treatment

Every neighbour that lands on $\partial\Omega$ is removed from the unknown list. If the stencil at interior index k has a Dirichlet neighbour with value g_{edge}, the term $a_{\text{face}}(U_k - g_{\text{edge}})/h^2$ contributes

$$b_k \mathrel{+}= \frac{a_{\text{face}}\, g_{\text{edge}}}{h^2} \tag{7}$$

to the right-hand side. For P5 ($u^* = e^x \sin(\pi y)$, $a = 1$) the west column $i = 0$ is the row set $k = 0, \ldots, n-1$; an example is

$$b[0] \mathrel{+}= \frac{e^0 \sin(\pi y_0)}{h^2}.$$

Replacing a Dirichlet row by e_k while leaving the corresponding column untouched would destroy symmetry [8]. Boundary nodes are therefore absent from A_h.

For the hand-checked small matrix ($n = 2$, $h = 1/3$, $a \equiv 1$, $k = 2i + j$), the four interior nodes form

$$A_h = \frac{1}{h^2}\begin{pmatrix} 4 & -1 & -1 & 0 \\ -1 & 4 & 0 & -1 \\ -1 & 0 & 4 & -1 \\ 0 & -1 & -1 & 4 \end{pmatrix} = \begin{pmatrix} 36 & -9 & -9 & 0 \\ -9 & 36 & 0 & -9 \\ -9 & 0 & 36 & -9 \\ 0 & -9 & -9 & 36 \end{pmatrix}, \qquad \operatorname{nnz}(A_h) = 12, \tag{8}$$

which matches `test_hand_matrix_n2` in `validation_tests.py`.

2.4 Discrete Structure, Stability, and Convergence

After Dirichlet elimination, every off-diagonal entry of A_h is non-positive and the row sums on interior stencils that do not touch the boundary are zero; boundary-adjacent rows are strictly diagonally dominant because a positive a_{face}/h^2 has been moved to b. With positive shared face coefficients, A_h is a symmetric M-matrix, hence SPD, and the discrete maximum principle holds: a discrete solution of $A_h U = 0$ with Dirichlet data in $[m, M]$ remains in $[m, M]$.

The continuous energy identity for (1) with homogeneous Dirichlet data is

$$\int_\Omega a|\nabla u|^2 \,\mathrm{d}x = \int_\Omega f u \,\mathrm{d}x, \tag{9}$$

so coercivity is controlled by $a_{\min} > 0$. On the grid, the shared-face assembly is the discrete analogue: $U^\top A_h U = \sum_{\text{faces}} a_{\text{face}}(\delta U)^2/h^2 \geq a_{\min}\,|U|^2_{H^1_h}$. Dirichlet elimination does not destroy that identity because boundary values enter b, not the quadratic form on interior unknowns. Consequently A_h remains SPD for every admissible a used in this report (smooth $a = 1+x$ and both harmonic/arithmetic layered averages).

Why the harmonic mean is forced at a jump. Across a vertical interface the exact flux $q = a u_x$ is continuous. On one grid edge of length h that straddles a_L and a_R in equal halves, integrating $u_x = q/a$ gives

$$U_R - U_L = q \cdot \tfrac{h}{2}\left(\tfrac{1}{a_L} + \tfrac{1}{a_R}\right) \quad \Rightarrow \quad q = \frac{2a_L a_R}{a_L + a_R}\,\frac{U_R - U_L}{h}. \tag{10}$$

So the face coefficient that preserves the continuous flux is exactly (6). The arithmetic mean $(a_L + a_R)/2$ is a different resistance and cannot match q unless $a_L = a_R$.

Convergence. On smooth data, consistency of order 2 plus the M-matrix / energy stability above imply convergence of order 2 in the discrete L^2 and L^∞ norms (Lax equivalence for linear FD schemes). Table 1 confirms the rate numerically: fitted slopes 2.00 with no deviation on the four meshes. On the layered jump the regularity assumption fails, so the theorem does not apply: harmonic averaging restores nodal exactness via (10), while arithmetic averaging drops to first order (Table 4).

Three numerical checks corroborate the derivation: $\|A - A^\top\|_{\max}/\|A\|_{\max} = 0$ on every assembled matrix in this report; the smallest algebraic eigenvalue of the P4 matrix at $n = 15$ is $\lambda_{\min} = 28.61 > 0$; the layered numerical solution stays in $[0, 1]$. Conditioning is separate: $\kappa_2(A_h)$ on P4 grows like h^{-2} (126.5, 529.7, 2.18×10^3 at $n = 15, 31, 63$), as expected for a second-order elliptic operator.

3 Implementation

3.1 Algorithm

Algorithm 1 Assembly of $A_h U = b$ and the direct solve

Require: interior count n, coefficient a, data f, g, face rule (half-point / harmonic / arithmetic)
Ensure: u_h, residual $\|b - A_h u_h\|_\infty$

1: $h \leftarrow 1/(n+1)$; index $k = in + j$
2: $b \leftarrow f$ sampled at interior nodes
3: **for** each interior node (i, j) **do**
4: **for** each of the four faces **do**
5: compute a shared $a_{\text{face}} > 0$
6: **if** the neighbour is interior **then**
7: add a_{face}/h^2 to the diagonal and $-a_{\text{face}}/h^2$ to that column
8: **else**
9: add a_{face}/h^2 to the diagonal and $a_{\text{face}} g_{\text{edge}}/h^2$ to b_k
10: **end if**
11: **end for**
12: **end for**
13: convert COO triplets to CSR; $u_h \leftarrow$ `spsolve`(A, b)
14: report $\|b - Au_h\|_\infty$

Lines 4–11 build one shared face at a time so that A_h inherits symmetry from the continuous flux form; the Dirichlet branch (lines 9–10) is exactly (7).

3.2 Code Structure

File	Purpose
`model.py`	$a(x, y)$, manufactured u^*, layered exact profile (3)
`assemble.py`	sparse face-averaged stencil, index map, interface flux
`solvers.py`	`spsolve`, residual-stopped CG, $\lambda_{\min}$ / κ via `eigsh`
`verify.py`	named-norm errors, four-mesh tables, layered diagnostics
`plots.py`	every report figure
`fem1d.py` / `fem2d_oracle.py`	independent 1D P1 and optional 2D P1 oracles
`test_problem_bank.py`	analytic bank (P1–P5), self-check of $f = -\nabla \cdot (a \nabla u^*)$
`validation_tests.py`	sign, boundary, order, nodal-exact, and M-matrix checks
`run_all.py`	regenerates every figure and prints the quoted numbers

3.3 Key Design Decisions

I assemble in COO using vectorised 2D index masks over the four stencil directions (no Python loops over grid nodes) and convert to CSR once. The physics (a, f, g) lives in `model.py`, so the same assembler runs P4, P5 and the layered jump. The default solve is sparse direct: at the meshes in this report ($n \le 254$, $N_{\text{unk}} \le 6.5 \times 10^4$, five nonzeros per row) a banded factorisation is cheaper than unpreconditioned CG, which I still run as a residual-stopped iterative check and as a contrast-sensitivity

probe. Because shared face coefficients make A_h symmetric positive definite, Krylov CG is preferred over stationary Jacobi or Gauss–Seidel (error contraction governed by $(\sqrt{\kappa}-1)/(\sqrt{\kappa}+1)$ rather than $(\kappa-1)/(\kappa+1)$). Cost measure (named before it is quoted): for CG, iteration count to a stated relative residual; for the direct solver, one sparse factorisation, with wall time only as a secondary observation.

4 Verification and Numerical Experiments

4.1 Verification: Method of Manufactured Solutions

The bank's finite-difference self-check on $f = -\nabla \cdot (a\nabla u^*)$ passes with worst relative error 5.8×10^{-6} (analytic data check, not a solver test) [8]. Sign and boundary gates come first: P2 ($u^* = x^2 - y^2$) has $L^\infty = 3.3 \times 10^{-16}$ at $n = 15$ (exact on this quadratic); P5 recovers the Tutorial 5 centre $u_h(0.5, 0.5) = 1.652$.

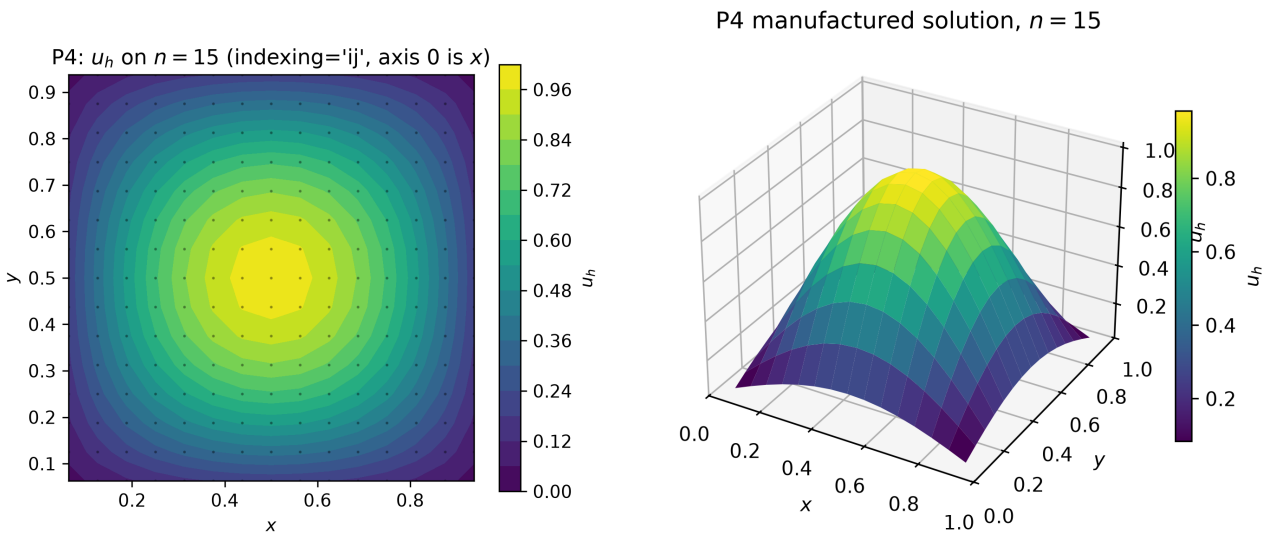
Table 1 is the P4 refinement study. Named norms: discrete L^2 given by $h\|u_h - u^*\|_2$, and $\|u_h - u^*\|_\infty$. Both fitted log-log slopes on $n = 15, 31, 63, 127$ are 2.00 (standard error of the fit $< 10^{-3}$). Figure 3 plots that table; Figure 1 shows the field and mesh; Figure 2 shows the sparsity pattern after elimination.

Table 1: P4 smooth manufactured solution, face coefficient at the half-point, `spsolve`. Discrete L^2 means $h\|u_h - u^*\|_2$.

n	h	L^2	L^∞	$\|b - Au_h\|_\infty$	p_{L^2}	p_∞
15	1/16	1.598×10^{-3}	3.205×10^{-3}	4.48×10^{-13}	—	—
31	1/32	3.988×10^{-4}	8.009×10^{-4}	3.00×10^{-12}	2.00	2.00
63	1/64	9.966×10^{-5}	2.003×10^{-4}	1.40×10^{-11}	2.00	2.00
127	1/128	2.491×10^{-5}	5.007×10^{-5}	6.92×10^{-11}	2.00	2.00

Table 2: P4 truncation $h\|b - A_h u^*\|_2$. Each refinement halves h and multiplies the truncation by $\approx 1/4$.

n	h	$h\|\tau_h\|_2$	$\|\tau_h\|_\infty$
15	1/16	4.978×10^{-2}	1.028×10^{-1}
31	1/32	1.248×10^{-2}	2.572×10^{-2}
63	1/64	3.122×10^{-3}	6.433×10^{-3}
127	1/128	7.808×10^{-4}	1.609×10^{-3}



(a) Filled contours with interior nodes ($n = 15$). **(b)** Surface plot of u_h.

Figure 1: P4 solution at $n = 15$. Mesh convention: `indexing='ij'`, axis 0 is x. Contours show the field the mesh actually resolved.

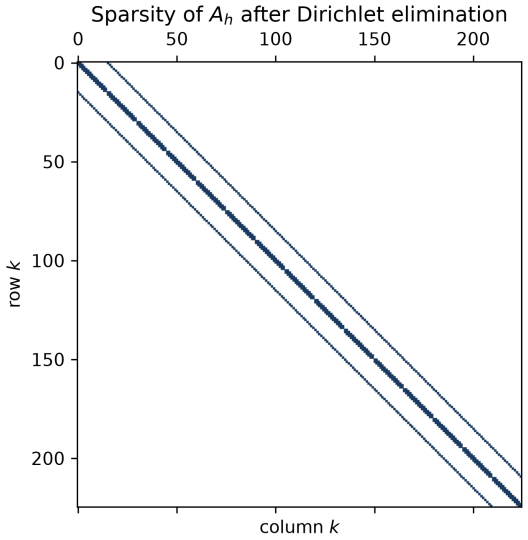


Figure 2: Spy plot of A_h at $n = 15$ after Dirichlet elimination (225×225). Five-point bandwidth matches the index map $k = in + j$.

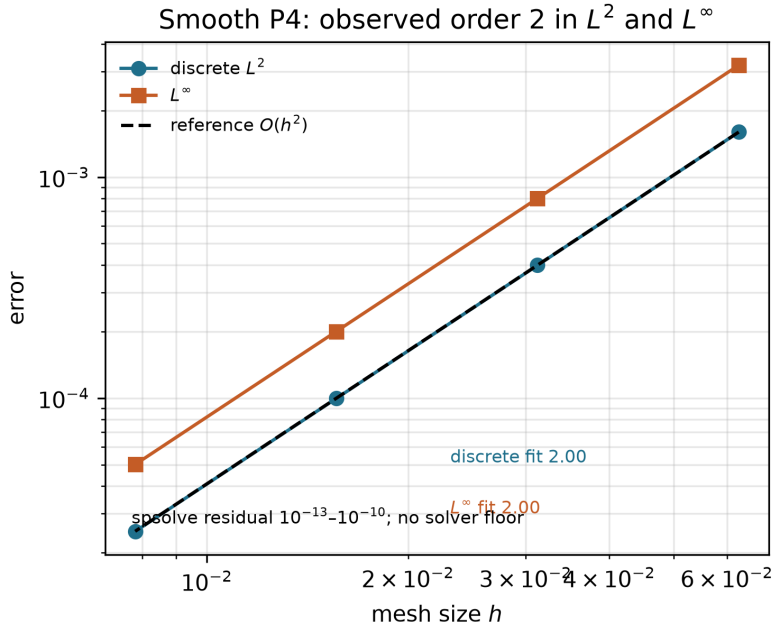


Figure 3: Log-log P4 errors versus h. Fitted discrete-L^2 order 2.001 ± 0.000, L^∞ order 2.000 ± 0.000, four meshes, reference slope 2. Confirms that the face-averaged stencil is second order on smooth data.

An independent 2D P1 FEM solve of the same P4 data (`scikit-fem`, homogeneous Dirichlet by `condense`) gives $L^2 = 5.18 \times 10^{-3}$ at $h_{\max} = 8.84 \times 10^{-2}$, residual 2×10^{-15}: an order-of-magnitude cross-check on a different mesh and method. The 1D P1 training problem $-u'' = \pi^2 \sin(\pi x)$ with nodal trapezoidal load fits discrete-L^2 order 2.00; the hand-checked element matrix is

$$K^{(e)} = \frac{1}{h}\begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix}, \qquad h = \tfrac{1}{4} \;\Rightarrow\; \begin{pmatrix} 4 & -4 \\ -4 & 4 \end{pmatrix}.$$

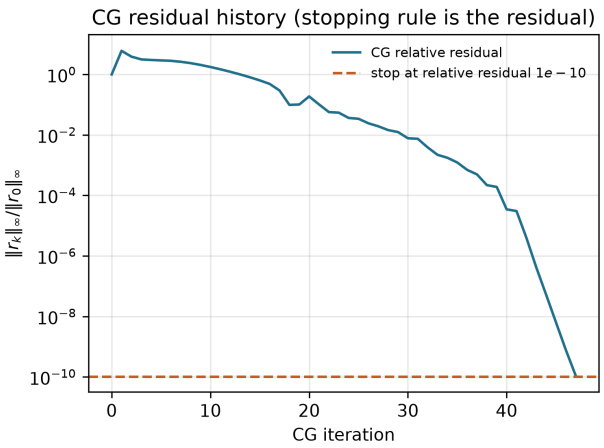
4.2 Separating Algebraic and Discretisation Error

`spsolve` has no tolerance parameter. I treat its solution as the exact discrete solution and report the backward residual. On P4 at $n = 31$, $\|b - Au_h\|_\infty = 3.00 \times 10^{-12}$ while the discrete L^2 error is 3.99×10^{-4}. A crude forward-error sketch $\|e\|_2 \lesssim \|r\|_2/\lambda_{\min}$ with $\lambda_{\min} \approx 28.7$ puts the algebraic error near 10^{-13}, four orders below the discretisation error.

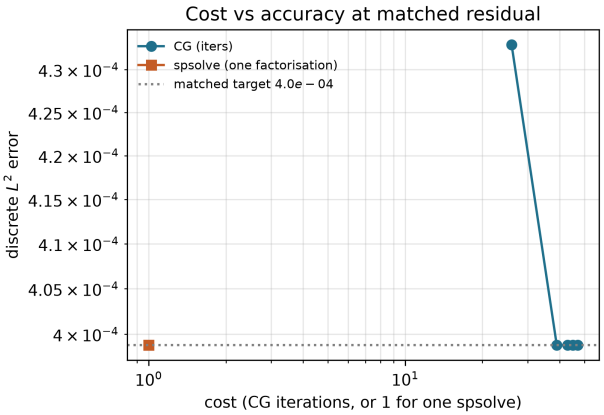
Unpreconditioned CG on the same matrix, stopped on the relative residual $\|r_k\|_\infty/\|r_0\|_\infty$, is the second demonstration (Figure 4). At relative tolerance 10^{-2} the discrete L^2 error is still 4.33×10^{-4}; from 10^{-4} onwards it has locked onto the direct value 3.988×10^{-4} (Table 3). Tightening further does not change the reported discretisation error, so the fitted order in Table 1 is the discretisation rate. I stop CG on the residual, never on $\|x^{k+1} - x^k\|$.

Table 3: P4, $n = 31$: CG versus `spsolve`. Cost for CG is iteration count; `spsolve` is one factorisation (0.002 s wall). The L^2 column is the quantity of interest.

Solve	iterations	$\|b - Au\|_\infty$	discrete L^2
CG, rel. 10^{-2}	26	7.5×10^{-1}	4.329×10^{-4}
CG, rel. 10^{-4}	39	5.9×10^{-3}	3.988×10^{-4}
CG, rel. 10^{-6}	43	1.3×10^{-5}	3.988×10^{-4}
CG, rel. 10^{-10}	47	3.5×10^{-9}	3.988×10^{-4}
`spsolve`	—	3.0×10^{-12}	3.988×10^{-4}



(a) CG relative residual vs. iteration.



(b) Discrete L^2 vs. CG iterations.

Figure 4: Left: residual history with stopping line at 10^{-10}. Right: discrete L^2 versus CG iterations; dashed line is the `spsolve` error. Once the residual is tight enough the error floor is discretisation. At $n = 31$ both solves finish in about 2 ms.

4.3 Depth Comparison: Harmonic versus Arithmetic Face Coefficients

This is the project's depth contribution beyond a smooth MMS: the same scheme, same grids, same Dirichlet data (3), same direct solver; only the face rule changes. The hypothesis under test is that "any positive face average that keeps A_h symmetric is an acceptable interface scheme." Table 4 and Figures 5–6 falsify that hypothesis and confirm (10).

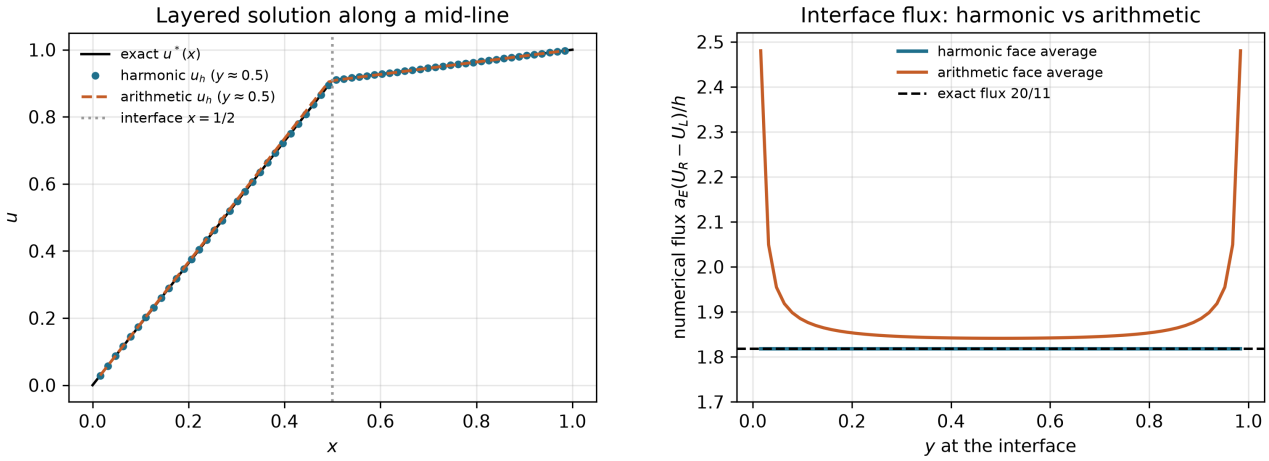
Harmonic averaging is nodally exact: $L^\infty \le 3 \times 10^{-14}$ on all four meshes, and the numerical flux $a_E(U_R - U_L)/h$ equals 20/11 to at most 5×10^{-13} (the pack's sanity check) [8]. An interface-fitted 1D P1 FEM with 32 elements (so $x = 1/2$ is a node) reproduces the same profile to $L^\infty = 6.7 \times 10^{-16}$. Two methods that share neither the stencil nor the mesh convention agree on the exact nodal values.

Arithmetic averaging still produces a symmetric M-matrix, but the wrong series resistance in (10). The flux jump stays ≈ 0.66 under refinement ($\mathcal{O}(1)$, not $\mathcal{O}(h)$). The L^∞ solution error is first-order ($\approx 1.94 \times 10^{-2}$ down to 2.38×10^{-3} as N runs through $31, 63, 127, 255$). Relative to the smooth-case energy (9), the flux error is the physically meaningful defect: arithmetic keeps a discrete energy form but with the wrong a_{face}, so L^2-looking plots can look stable while transport is wrong.

On every harmonic and arithmetic run the discrete solution stays in $[0, 1]$: both averages inherit the M-matrix / DMP. Monotonicity alone does not select the correct face rule. Value of the comparison: it separates algebraic structure (symmetry, SPD) from interface modelling (flux continuity), which is the core numerical lesson of Direction 3.

Table 4: Layered $a = 1 \,|\, 10$, odd interval counts so $x = 1/2$ bisects a grid edge. Flux jump is $\max_j |a_E(U_R - U_L)/h - 20/11|$; p_∞ is the observed arithmetic L^∞ order.

n	N	harm. L^∞	harm. flux jump	arith. L^∞	arith. p_∞	arith. flux jump
30	31	2.4×10^{-15}	3.9×10^{-14}	1.94×10^{-2}	—	0.663
62	63	7.9×10^{-15}	1.1×10^{-13}	9.60×10^{-3}	1.01	0.662
126	127	2.6×10^{-14}	1.8×10^{-13}	4.78×10^{-3}	1.01	0.662
254	255	3.5×10^{-14}	4.6×10^{-13}	2.38×10^{-3}	1.00	0.662



(a) Profile along $y \approx 0.5$.

(b) Interface flux versus y.

Figure 5: Left: harmonic u_h (blue circles) and arithmetic u_h (orange dashed) against the exact piecewise-linear profile; the kink at $x = 1/2$ sits between columns. Right: interface flux versus y. The dashed line is the exact flux $20/11$; harmonic sits on it, while arithmetic is a different constant that does not shrink with h.

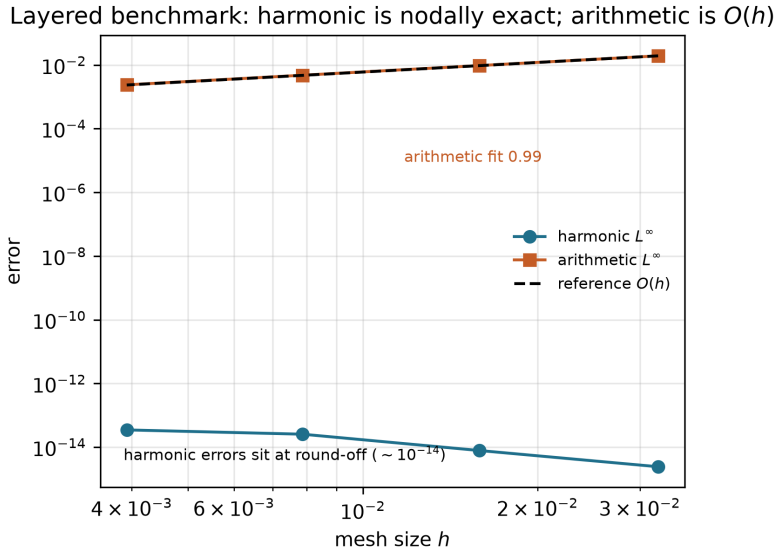


Figure 6: Layered L^∞ versus h. Harmonic errors sit at round-off ($\sim 10^{-14}$). Arithmetic tracks the $O(h)$ reference line (fitted slope 0.99). Same data, two face rules, named norm: the depth comparison.

4.4 Main Results: Contrast, Conditioning, Cost

Unpreconditioned CG iteration counts on the layered harmonic matrix at $n = 30$, relative residual 10^{-10}, grow with the contrast a_R/a_L: 125, 227, 401, 615 iterations at contrasts $2, 10, 50, 200$ (Figure 7). Coefficient contrast worsens conditioning but preserves the discrete maximum principle. On smooth P4, $\kappa_2(A_h)$ tracks h^{-2} as already quoted.

Cost–accuracy at a matched target. Cost measure (stated before quoting): CG iteration count to a relative residual, versus one sparse factorisation for `spsolve`. Accuracy target (frozen across the comparison): discrete L^2 error equal to the `spsolve` value 3.988×10^{-4} on P4 at $n = 31$. Table 3 and Figure 4 (right) show that CG first hits this target at relative residual 10^{-4} (39 iterations, ≈ 2 ms wall); `spsolve` reaches the same L^2 in one factorisation (≈ 2 ms). At matched discretisation accuracy there is no runtime crossover for $n \leq 31$.

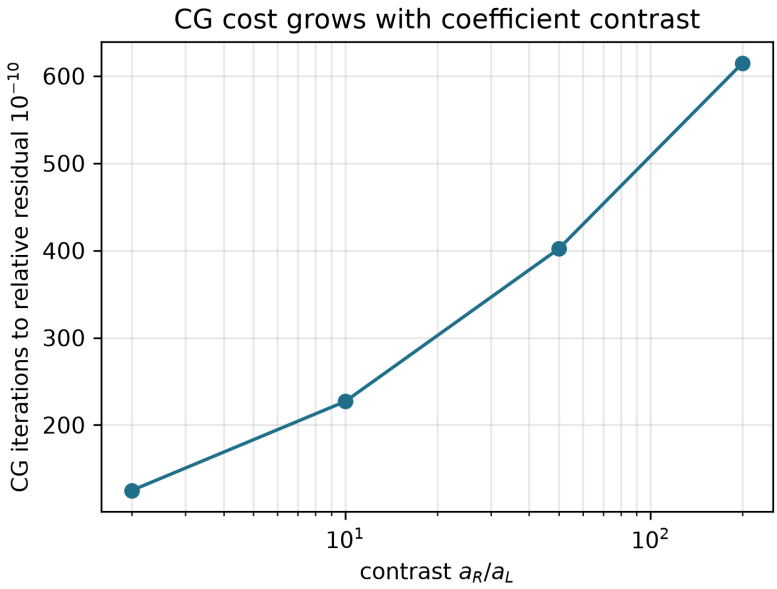


Figure 7: CG iterations versus conductivity contrast on the layered problem, $n = 30$, residual stop 10^{-10}. Cost measure: iterations. Contrast raises iterative cost without breaking the discrete maximum principle.

4.5 Limitations and Robustness

The second-order claim applies to smooth a with the half-point face rule. Across a jump, harmonic averaging on an aligned grid is exact for this piecewise-linear u^*. Arithmetic averaging remains an M-matrix and still looks smooth on a colour plot, yet its flux is wrong. Unpreconditioned CG will not be competitive on finer 2D meshes; AMG-CG (Tutorial 6) is the natural next solver.

5 Conclusions

- A conservative face-averaged five-point scheme for $-\nabla \cdot (a\nabla u)$ with Dirichlet data moved into b was assembled and verified on the analytic bank before the jump.
- On smooth P4, discrete L^2 and L^∞ both fit order 2.00 over four meshes; $\|A - A^\top\|_{\max} = 0$ and $\lambda_{\min} > 0$. The reported rate survives a CG tolerance sweep.
- On the layered jump, harmonic face coefficients are nodally exact and keep the flux at 20/11; arithmetic coefficients give first-order L^∞ error and an $\mathcal{O}(1)$ flux jump. Symmetry is necessary but not sufficient for correct interface physics.
- Cost is CG iterations or one sparse factorisation. At $n = 31$, CG and `spsolve` agree in L^2 once the relative residual is 10^{-4}. Next step: AMG-preconditioned CG on finer meshes and non-aligned interfaces.

References

[1] W. Hackbusch, *Elliptic Differential Equations: Theory and Numerical Treatment*, 2nd ed. Berlin: Springer, 2017.

[2] R. J. LeVeque, *Finite Difference Methods for Ordinary and Partial Differential Equations*. Philadelphia: SIAM, 2007.

[3] R. Hiptmair, *Numerical Methods for PDEs*, lecture notes, ETH Zürich, 2025.

[4] A. Quarteroni, *Numerical Models for Differential Problems*, 3rd ed. Cham: Springer, 2017.

[5] P. Virtanen et al., "SciPy 1.0: fundamental algorithms for scientific computing in Python," *Nature Methods*, vol. 17, pp. 261–272, 2020.

[6] C. R. Harris et al., "Array programming with NumPy," *Nature*, vol. 585, pp. 357–362, 2020.

[7] T. Gustafsson and G. D. McBain, "scikit-fem: a Python package for finite element assembly," *Journal of Open Source Software*, vol. 5, no. 52, 2020.

[8] Course problem pack, "Variable diffusion," Project 2 Direction 3, Numerical Analysis of ODEs and PDEs, 2026.

[9] Course brief, "Numerical Solution of Elliptic Boundary Value Problems," Project 2, 2026.

A AI Transparency Log

I used Grok (via Cursor) to help write the Python code and to polish the English of this report.

Item	Claim / assumption	Risk / falsifier	Independent check & decision
1	Conservative stencil (5) with Dirichlet shift $b_k \mathrel{+}= a_{\text{face}} g/h^2$	Sign error in boundary shift or unsymmetric boundary row	Hand matrix (8) (nnz = 12), P2 $L^\infty = 4.4 \times 10^{-16}$, P5 centre 1.652; accepted
2	Layered exact profile (3) and harmonic vs. arithmetic face average	Using a discontinuous OCR formula or treating symmetry alone as proof of flux accuracy	Continuity of u and $au_x = 20/11$ checked analytically; Table 4 shows harmonic $L^\infty \sim 10^{-14}$ vs. arithmetic $O(h)$

Ownership and library interfaces: Self-implemented in `code/assemble.py`, `model.py`, `verify.py`, `fem1d.py`. Library calls: `scipy.sparse.coo_matrix` / `spsolve` (direct CSR solve), `scipy.sparse.linalg.cg` (residual-stopped CG), `eigsh` ($\lambda_{\min}, \kappa_2$), and `scikit-fem` (2D P1 cross-check).

B Individual Contribution Statement

Individual Contribution Statement — Project 2 (PDE BVP)
Name: Chuyang Wang Student number: 2363591
Role this project: sole author (implementation, weekly tasks, report).

The specific things I did this project:

1. Derived the conservative stencil (5), the Dirichlet shift (7), and the $n = 2$ hand matrix (8); implemented vectorised COO assembly in `code/assemble.py` with the index map $k = in + j$ (cf. Appendix A, Item 1).

2. Completed the weekly individual challenges through Week 8: Week 5 five-point P5 centre $u_h(0.5, 0.5) = 1.652$; Week 6 four-mesh slope (2.00 in discrete L^2 and L^∞); Week 7 residual/tolerance check separating algebraic from discretisation error (Table 3); Week 8 feedback-to-evidence improvement on harmonic versus arithmetic interface flux (Table 4, Figures 5–6).

3. Re-derived the layered profile (3), implemented harmonic versus arithmetic face averages, and produced Table 4 and Figures 5–6 (cf. Appendix A, Item 2).

4. Separated algebraic from discretisation error by a CG residual sweep (Table 3); measured CG iterations against contrast (Figure 7).

5. Wrote `validation_tests.py` and `run_all.py`, and wrote the full LaTeX report including the AI log and this ICS.

What I would do differently: freeze the layered exact formula on paper before reading the OCR'd pack, instead of debugging a "failed" nodal-exact test against a discontinuous u^*.
Anything I should have asked for help with earlier: whether odd N (interface between columns) versus even N (interface on a node) is required for the FDM sanity check, before running the first layered table on $n = 15$.

C Supplementary data

P4 structure at $n = 15$: $\|A - A^\top\|_{\max}/\|A\|_{\max} = 0$, $\lambda_{\min} = 28.61$, $\kappa_2 \approx 1.27 \times 10^2$.
FEM 1D element matrix for $h = 1/4$: $\left(\begin{smallmatrix} 4 & -4 \\ -4 & 4 \end{smallmatrix}\right)$.
All numbers in the tables are reproduced by `python code/run_all.py`.